pymodalhammer — User Manual

Impact-hammer modal testing with NI-DAQmx · version 0.1.0b1 (first beta)

The pymodalhammer authors

November 2026

Contents

1 About pymodalhammer **3**
1.1 What it does . . . 3
1.2 Status of this version . . . 3
1.3 Licence . . . 4

2 Installation **4**
2.1 Requirements . . . 4
2.2 Installing . . . 4
2.3 First start and simulation mode . . . 4

3 Quick start **4**

4 The main window **5**

5 Channel Setup **6**

6 Acquisition Setup **7**
6.1 Record . . . 7
6.2 Trigger (hit detection) . . . 8
6.3 Windows . . . 9
6.4 Averaging and acceptance . . . 11
6.5 Hit checks — what to do with a bad hit . . . 12
6.6 External trigger (e.g. a high-speed camera) . . . 13

7 Probe mode **14**

8 Measurement **16**
8.1 Save options . . . 16
8.2 Measurement point . . . 16
8.3 Running a point . . . 17
8.4 The hit list . . . 17
8.5 Manual review and “Ask me” . . . 18
8.6 Time domain, spectra and FRF . . . 18
8.7 Roving-hammer workflow, step by step . . . 20
8.8 Time data only (hammer channel alone) . . . 20

9 Data Viewer **21**
9.1 Hits and re-processing . . . 21
9.2 FRF, time data and spectra . . . 22
9.3 Compare runs . . . 24

10 Saved files **26**
10.1 UFF 58 records . . . 27
10.2 TDMS hit file . . . 27
10.3 CSV logs . . . 27
10.4 Reading the files in Python . . . 28

11 How it works **28**
11.1 Hit detection . . . 28
11.2 Spectra and FRF estimators . . . 29
11.3 Checks in detail . . . 29

12 Hardware notes **29**

13 Troubleshooting **30**

14 Testing and development **30**

15 Reporting problems **31**

16 Licence **31**

1 About pymodalhammer

pymodalhammer is a desktop program for **impact-hammer modal testing** with National Instruments data-acquisition hardware (CompactDAQ chassis such as the cDAQ-9174 with IEPE input modules such as the NI 9234). You strike a structure with an instrumented hammer; the program detects every hit, checks it, averages the good ones and computes frequency response functions (FRFs) between the hammer force and every response sensor. Results are saved in standard formats (UFF 58 for FRFs, TDMS for time data) that open in modal-analysis software, DIAdem, LabVIEW or Python.

1.1 What it does

- **Continuous acquisition with a pre-trigger.** The inputs run without gaps; every hit on the hammer channel is captured with a part of the signal from just before the impact. No hit is missed between records.
- **Quality checks on every hit**: overload, double hit (hammer bounce), hammer force outside a range, and responses that have not died out by the end of the record. For each check you choose what happens: reject, ask you, accept but flag, or ignore.
- **Automatic or manual acceptance.** In manual mode every hit waits for your Accept / Reject.
- **Full control over earlier hits.** Any hit can be deleted or put back into the average at any time; the FRF updates instantly.
- **Force and exponential windows**, with the damping added by the exponential window shown.
- **H1, H2, Hv and coherence**, averaged over the hits you keep.
- **Probe mode**: tap the structure and see the level of every channel, the hammer's frequency content and suggestions for the trigger level and input ranges, without measuring.
- **Roving-hammer workflow**: the hammer position is the measurement point; after saving, the point number advances automatically, optionally without stopping the acquisition.
- **External trigger output** (e.g. for a high-speed camera): a pulse at the start, on every hit, on every accepted hit, on the first hit of a point or when a point is complete, or a frame clock locked to the sample clock. Every pulse is logged with its exact sample index.
- **Data Viewer**: FRFs, time data and spectra of saved points; re-processing with other hits or other windows; comparison of several points (FRF, amplitude spectrum and time data).
- **Time data only**: with just the hammer channel enabled (no responses), the hits are recorded and saved as time data, with the same trigger, checks and external trigger output, but without FRFs.
- **Simulation mode**: a simulated hammer and structure, so everything can be tried without hardware.

1.2 Status of this version

This is **version 0.1.0b1, the first beta**. All features are complete. They are covered by automated tests in simulation mode and by tests of the NI-DAQmx code against a mock driver that reproduces the limits of real modules (e.g. the NI 9260). Behaviour on your own hardware has to be confirmed on your system: check results against a known reference before relying on them, and report problems (see 15).

1.3 Licence

pymodalhammer is free software, released under the **GNU General Public License, version 3 or (at your option) any later version** (GPL-3.0-or-later). You may use, study, share and modify it; if you distribute modified versions, they must be released under the same licence. The program comes **without any warranty**. The full text is in the file `LICENSE`.

2 Installation

2.1 Requirements

Item	Notes
Python	3.9 or newer
NI-DAQmx driver	from ni.com, only needed with hardware (Linux and Windows are supported by NI)
Python packages	PySide6, pyqtgraph, numpy, scipy, nidaqmx, pyuff, npTDMS (installed automatically)
Hardware (optional)	an NI CompactDAQ chassis with IEPE input modules; optionally an AO or DO module for the external trigger

2.2 Installing

From the project folder (the one that contains `pyproject.toml`):

```
python -m pip install .            # installs pymodalhammer and its dependencies
pymodalhammer                      # starts the program
```

Alternatively, without installing:

```
python -m pip install -r requirements.txt
python -m pymodalhammer            # run from the project folder
```

A virtual environment (`python -m venv .venv`) keeps the packages separate from the system's.

2.3 First start and simulation mode

At start-up the program looks for NI devices. If it finds none (or the driver is not installed), it starts in **Simulation mode**: the status bar shows `SIMULATION MODE` in orange. A simulated hammer then strikes a simulated structure with three modes (at about 123, 348 and 696 Hz with the default sampling rate) every 1.6–3.2 s. About one hit in ten is a double hit and the occasional hard hit overloads the hammer channel, so every feature can be tried. Use **Hardware → Simulation mode** to switch, **Hardware → Refresh NI devices** (F5) after connecting the chassis, and **Hardware → Simulator: random hits** to make the simulated hammer strike only when you press **Simulate a hit**.

3 Quick start

The fastest way to learn the program is to do these steps once in simulation mode.

1. **Channel Setup**: check that the hammer row is ticked as *Reference* and that the accelerometers are enabled (see 5).
2. **Acquisition Setup → Probe mode**: press *Start probe* and tap the structure (or wait for the simulator). Check the level meters and the hammer spectrum; press *Use suggested trigger level*. Press *Stop* (see 7).

3. **Acquisition Setup**: check the record length, the windows and the number of averages (see 6).
4. **Measurement**: choose the output folder and the hammer node of the first point.
5. Press **Start**. Wait until the status reads `ARMED`, then hit the structure. Each hit appears in the list; the plots and the FRF update.
6. When the point has its averages the acquisition stops. Look through the hits, delete any you do not like, and press **Save point**. The node number advances by one.
7. Move the hammer to the next point and repeat from step 5.
8. **Data Viewer**: look at the saved points and compare them (see 9).

4 The main window

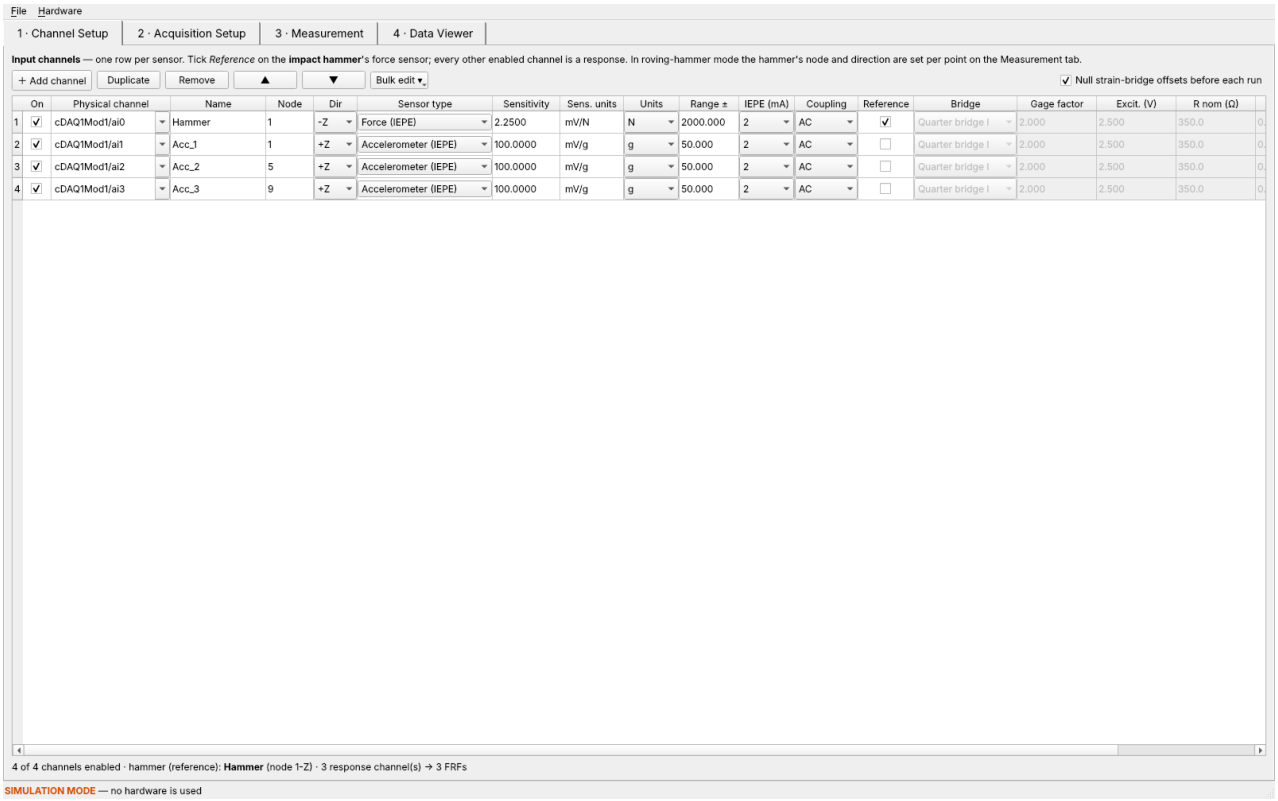


Figure 1: The main window with its four tabs. The status bar at the bottom shows whether the program works with hardware or in simulation.

The work is organised in four tabs, used from left to right:

Tab	Purpose
1 · Channel Setup	which sensors are connected where, their sensitivities, ranges, nodes and directions
2 · Acquisition Setup	record, trigger, windows, averaging, hit checks, external trigger; Probe mode
3 · Measurement	what to save, the measurement point, the hits and the live results
4 · Data Viewer	saved points: FRFs, time data, spectra, re-processing and comparison

Menus

- **File**: *New setup* (Ctrl+N), *Open setup...* (Ctrl+O), *Save setup* (Ctrl+S), *Save setup as...*, *Quit* (Ctrl+Q). A setup is a JSON file with every setting of every tab. The last setup is restored automatically when the program starts.
- **Hardware**: *Refresh NI devices* (F5), *Simulation mode*, *Simulator: random hits*.

While an acquisition runs, the settings that must not change are locked; they unlock when it stops. Every tab scrolls if the screen is smaller than its content, so the window can always be maximised.

5 Channel Setup

One row per sensor. The table is the same as in the companion shaker program.

Column	Meaning
On	include the channel
Physical channel	DAQmx name, e.g. `cDAQ1Mod1/ai0` (pick from the list or type)
Name	used in plots, files and to match channels between runs; keep names unique
Node, Dir	geometry node and direction (±X/Y/Z) written into the UFF files
Sensor type	accelerometer, force, microphone, pressure (IEPE), strain gauge, voltage or custom
Sensitivity	from the calibration sheet, in the units shown in *Sens. units* (e.g. mV/N for the hammer)
Units, Range ±	engineering unit and the largest expected level; the range also sets the overload check
IEPE (mA), Coupling	IEPE excitation current (0 = off) and coupling (AC for IEPE sensors)
Reference	**tick this on the hammer.** FRFs are response / hammer force
Bridge …	strain-gauge settings (bridge type, gage factor, excitation, resistance, Poisson's ratio)

Use + *Add channel*, *Duplicate*, *Remove* and the arrows to edit the list. *Bulk edit* numbers nodes, copies sensor settings to several rows, or adds every channel of a module. The line under the table summarises the setup (number of channels, the hammer, the number of FRFs).

> **Tip.** Typical impact hammers have a sensitivity of 1–25 mV/N. With an NI 9234 (±5 V input) a hammer of 2.25 mV/N can measure up to about 2200 N. Probe mode suggests sensible ranges.

In a **roving-hammer** test the hammer's node and direction in this table are not used: they are set per measurement point on the Measurement tab (see 8.2).

If the hammer is the only enabled channel, the measurement records **time data only** (see 8.8); the summary line says so.

6 Acquisition Setup

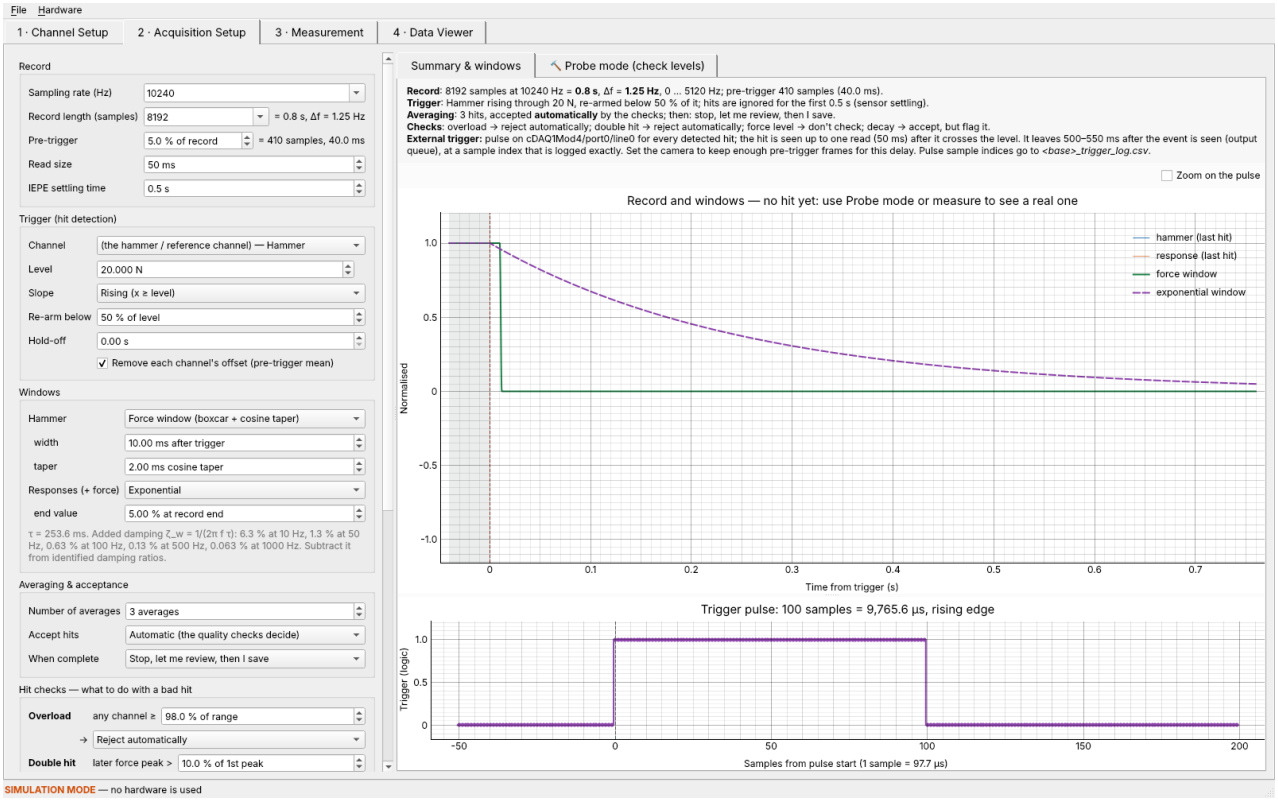


Figure 2: Acquisition Setup. Left: all settings (the column scrolls). Right: a summary in plain words, a preview of the record and its windows, and the external trigger pulse.

The summary on the right is rewritten whenever a setting changes. Read it before measuring: it states the record length, the frequency resolution, when the trigger fires, what happens to bad hits and when external-trigger pulses leave.

6.1 Record

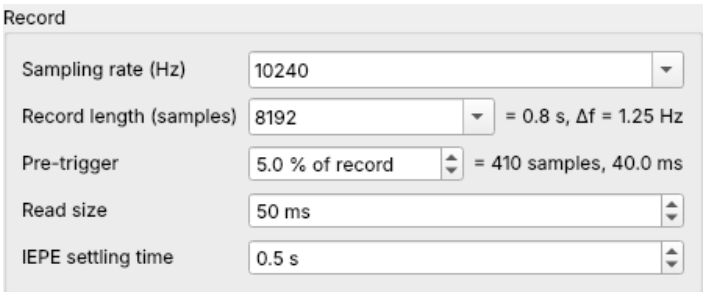


Figure 3: Record settings.

- **Sampling rate**: the analysis goes up to half of it (the Nyquist frequency). Delta-sigma modules such as the NI 9234 only allow rates of 51.2 kHz / n (e.g. 10 240, 12 800, 25 600 Hz); if another rate is entered, DAQmx uses the nearest allowed one and the program uses, logs and saves the rate actually used.
- **Record length** N (samples): the record time is $T = N/f_s$ and the frequency resolution $\Delta f = f_s/N$. Choose T long enough for the response to die out (or nearly so — the exponential window does the rest).

- **Pre-trigger**: the part of the record before the trigger point, in % of the record. A few percent is enough to see the start of the pulse and to estimate the sensor offsets.
- **Read size**: the acquisition is read in pieces of this length (default 50 ms). It sets how quickly a hit is noticed and how often the screen updates. Keep it ≥ 20 ms over USB.
- **IEPE settling time**: hits are ignored for this long after Start, while IEPE sensors and AC coupling settle.

Goal	Do
finer frequency resolution	longer record (more samples, or a lower rate)
higher frequency range	higher sampling rate (and a harder hammer tip)
lightly damped structure	longer record and/or a stronger exponential window

6.2 Trigger (hit detection)

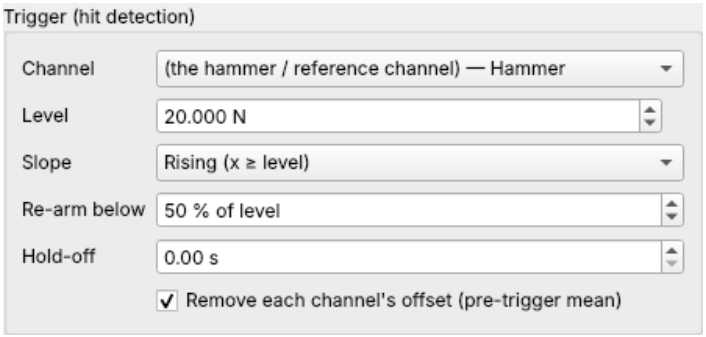


Figure 4: Trigger settings.

- **Channel**: normally the hammer (the reference channel); any channel can be used.
- **Level** and **Slope**: a hit is detected when the signal crosses the level — *rising* ($x \geq$ level), *falling* ($x \leq -$level) or *either* ($|x| \geq$ level). The first sample past the level is the **trigger point**, time $t = 0$ in all plots and files. Set the level well above the noise and well below a light tap (Probe mode suggests 10 % of a typical hit).
- **Re-arm below**: after a record, the trigger only re-arms once the signal has dropped below this fraction of the level (hysteresis). A ringing hammer therefore cannot trigger twice.
- **Hold-off**: an optional extra dead time after each record.
- **Remove each channel's offset**: subtracts the mean of the pre-trigger part of every channel.

Because the acquisition is continuous and buffered, a record never overlaps the previous one and the pre-trigger is always available, however soon after the previous record a hit comes.

6.3 Windows

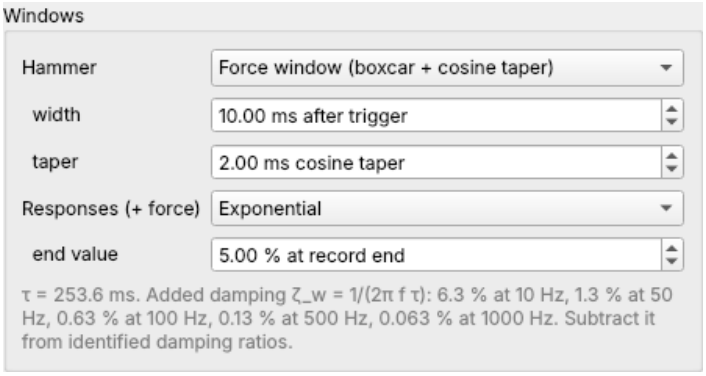


Figure 5: Window settings. The line below them gives the time constant of the exponential window and the damping it adds.

- **Hammer: force window**. The hammer signal is kept from the start of the record to *width* after the trigger point, then brought to zero by a half-cosine over *taper*. This removes the noise the hammer channel picks up after the impact. The width must cover the whole pulse (typically 1–5 ms) — check it on a real hit (figure below).
- **Responses: exponential window**. From the trigger point on, every channel is multiplied by $e^{-t/\tau}$, chosen so that it has fallen to *end value* at the end of the record. Responses that would still ring at the end of the record are forced to decay, which avoids leakage. The window is applied to the hammer too, so the FRF stays consistent.

The exponential window adds damping to every mode. The added damping ratio at frequency f is

$$\zeta_w = \frac{1}{2\pi f \tau}$$

and the panel lists it at a few frequencies. **Subtract it from damping ratios identified from the FRFs.** The time constant τ is saved with every point. Use the mildest window that removes the leakage (an end value of 1–10 % is typical); with a long enough record use none.

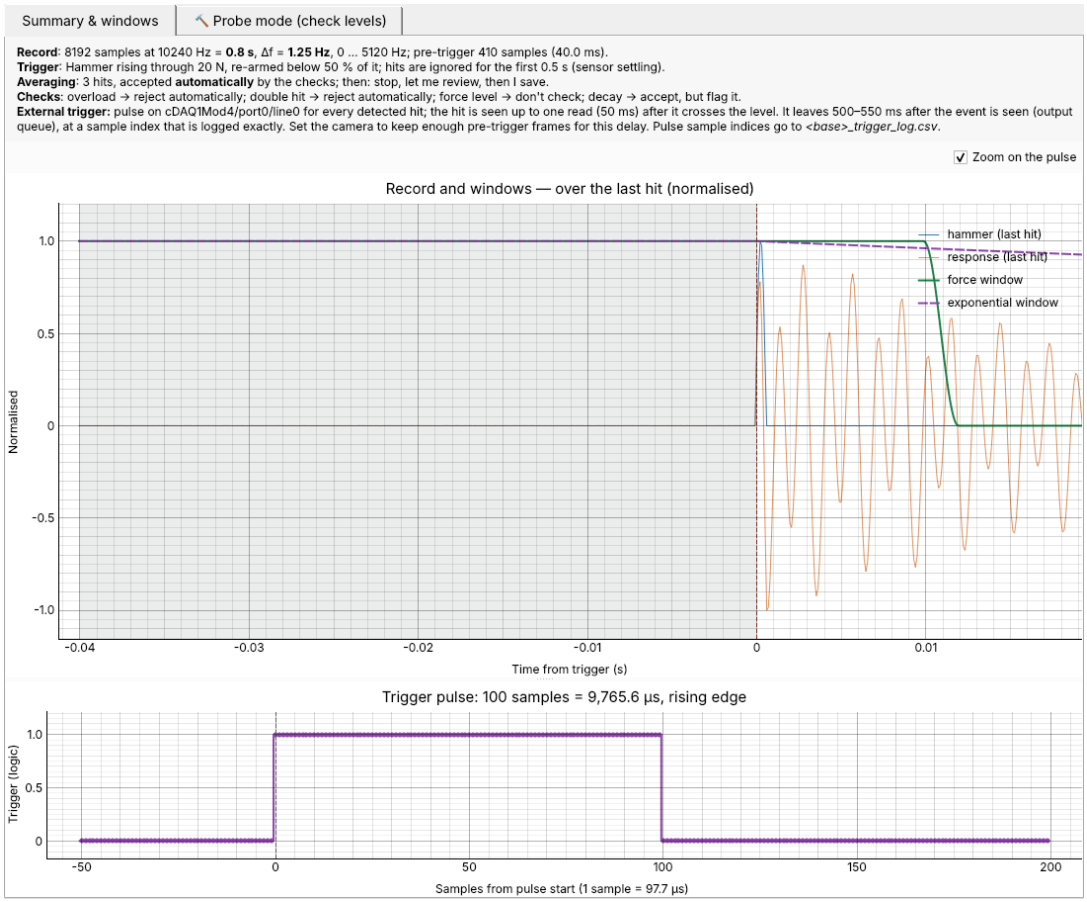


Figure 6: The window preview drawn over a real hit (from Probe mode or a measurement), zoomed on the impact: hammer (blue), first response (orange), force window (green) and exponential window (purple, dashed). Below: the external trigger pulse.

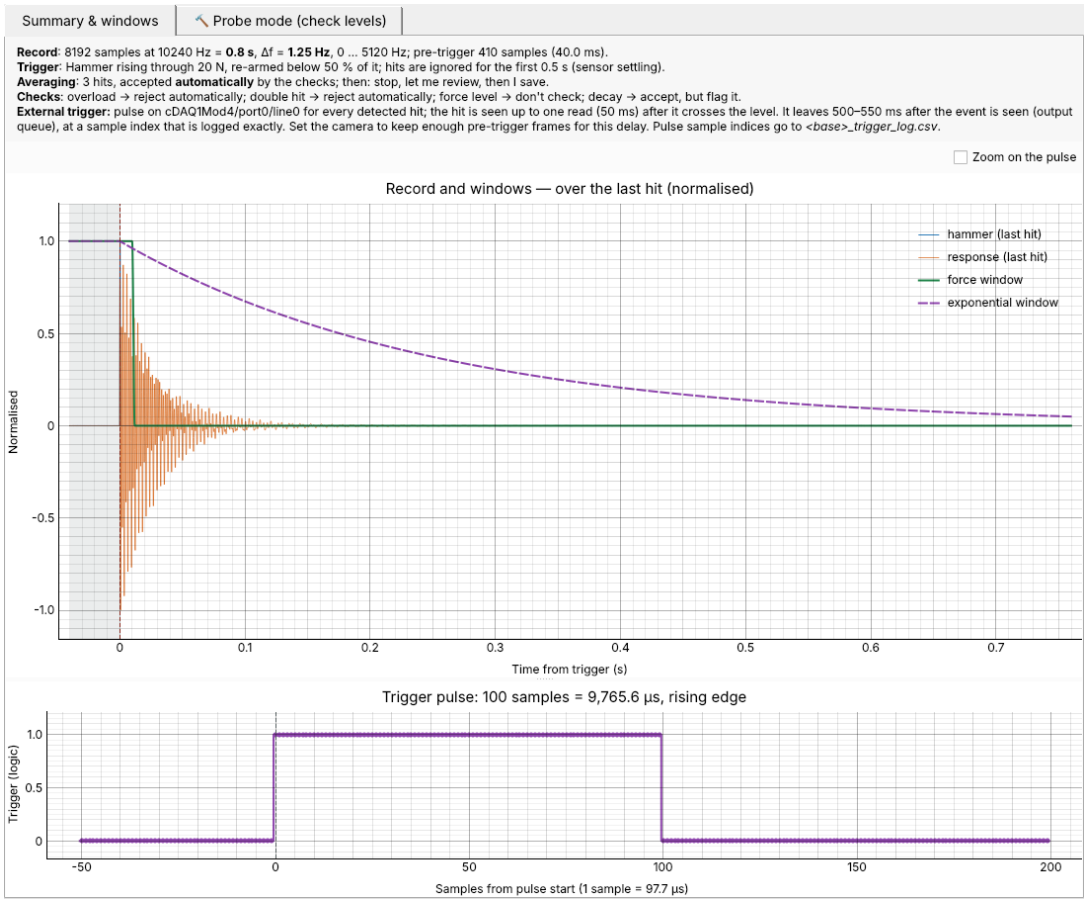


Figure 7: The same preview over the whole record: the response decays inside the exponential window.

6.4 Averaging and acceptance

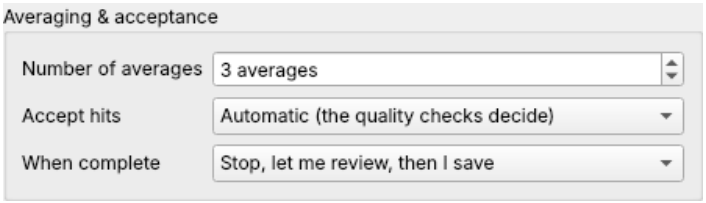


Figure 8: Averaging and acceptance.

- **Number of averages**: how many good hits a point needs (typically 3–10).
- **Accept hits**:
  - *Automatic*: the hit checks decide (next section).
  - *Manual*: every hit waits for you to press **Accept (A)** or **Reject (R)**; the check results are shown as advice. The trigger stays disarmed until you decide.
- **When complete**:
  - *Stop, let me review, then I save* (default): the acquisition stops; you look through the hits and press *Save point*.
  - *Stop and save automatically*.
  - *Save, then re-arm for the next point (roving hammer)*: the point is saved, the hammer node advances and the trigger re-arms **without stopping the acquisition** — the

fastest way through a grid of points.

6.5 Hit checks — what to do with a bad hit

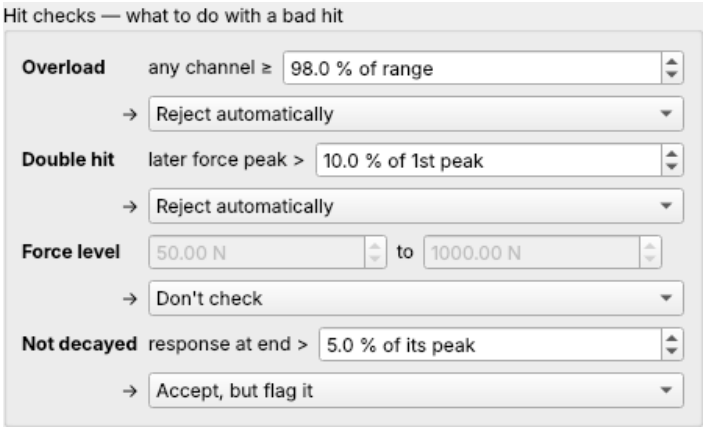


Figure 9: Hit checks: a threshold and an action per check.

Check	A hit fails when	Default action
Overload	any channel reaches the given % of its range	reject
Double hit	after the main pulse has ended (force below 5 % of its peak), the force rises again above the given % of the peak	reject
Force level	the peak hammer force is outside the given minimum … maximum	don’t check
Not decayed	a response, over the last 5 % of the record, is still above the given % of its peak (RMS·√2)	accept, but flag it

The actions:

Action	Effect
Reject automatically	the hit is not averaged; it stays in the list (red) and the trigger re-arms at once
Ask me	the hit waits for Accept / Reject (orange bar on the Measurement tab); the trigger is disarmed
Accept, but flag it	the hit is averaged but marked (orange) so you can review and delete it later
Don’t check	the check is switched off

If several checks fail, the strictest action wins (reject before ask before flag).

6.6 External trigger (e.g. a high-speed camera)

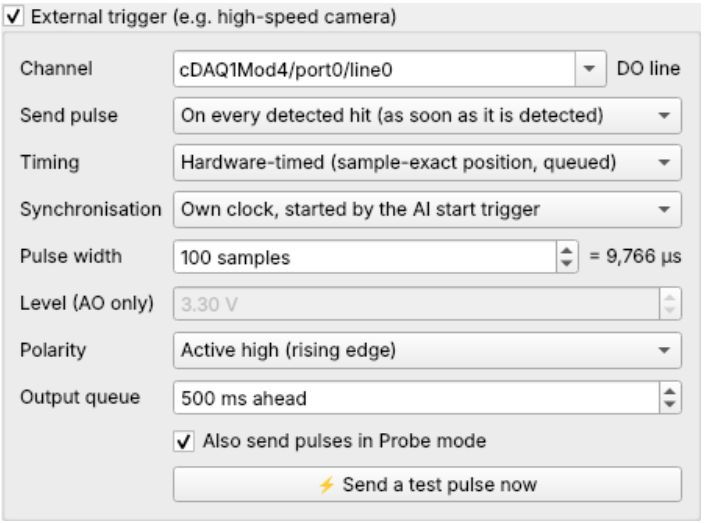


Figure 10: External trigger settings.

Tick the box to send a pulse to another device. Typical uses: start or mark the recording of a high-speed camera, tell a robot or an automatic hammer that a point is done, or mark hits on a second acquisition system.

Channel: an analog output (`…/aoN`) or a digital line (`…/portN/lineM`); for frame sync, a PFI terminal (`/cDAQ1Mod4/PFI0`). The label next to it shows how the channel was recognised.

Send pulse — when:

Mode	The pulse is sent
Once, when the acquisition starts	at sample 0 (exactly)
On every detected hit	as soon as a hit crosses the trigger level (before the record is complete and before any check)
On every accepted hit	when the record is complete and passed the checks — or when you press Accept
On the first hit of each measurement point	for the first hit after a point is started
When the averages of a point are complete	when the last average has been accepted
Frame sync	the input sample clock itself is routed to a PFI pin: one pulse per sample, for the whole run

Timing:

- *Hardware-timed* (default, recommended). The trigger line is driven by a continuous output task that runs in step with the inputs. The output is prepared ahead of time (the **output queue**, default 500 ms) and a pulse is written into the next part to go out. So **a pulse leaves about one output queue after the event**, at a sample index that is known exactly and logged. The output queue is also the slack the computer has: if the program is held up for longer, the output runs dry and the run stops with an “underflow” message. Increase the queue to 1000–2000 ms on a busy computer.
- *On-demand*. The line is switched by the program as soon as the event is seen (a few ms over USB). The logged sample index is approximate. Needs an output that allows software-timed writes: digital lines and most AO modules, but **not** delta-sigma AO modules such as the NI 9260.

Synchronisation (hardware-timed only):

- *Own clock, started by the AI start trigger* (default): works with every module combination.
- *Clocked by the AI sample clock*: exact lockstep, but many module combinations refuse it; the program then falls back to the first option by itself and says so in the log.

Pulse width (samples, with the time shown), **Level** (analog outputs only; 3.3 V is enough for TTL inputs, which read ≥ 2 V as high), and **Polarity** (active high = rising edge, active low = falling edge). *Also send pulses in Probe mode* lets you check the wiring while probing. **Send a test pulse now** sends one pulse when nothing is running.

> **Using a camera with a hit-based pulse.** A hit can only be detected after it has happened, so a pulse "on every hit" always comes *after* the impact: about 0.5 s later with the default output queue. Use a camera with a ring (pre-trigger) buffer and set it to keep at least that much before the trigger (end-trigger or centre-trigger mode). The trigger log gives every pulse's sample index, so camera frames can be aligned with the hit records afterwards: frame offset = (pulse sample − hit trigger sample) × fps / fs.
>
> **Frame sync** removes all of this: the camera, in external-sync mode, exposes one frame per input sample, so frame f is sample f, and each hit's trigger sample (in the hit log) is its frame number.

7 Probe mode

Probe mode answers "what am I going to measure?" before you measure. It is on the right of the Acquisition Setup tab, under **Probe mode (check levels)**. Nothing is averaged or saved.

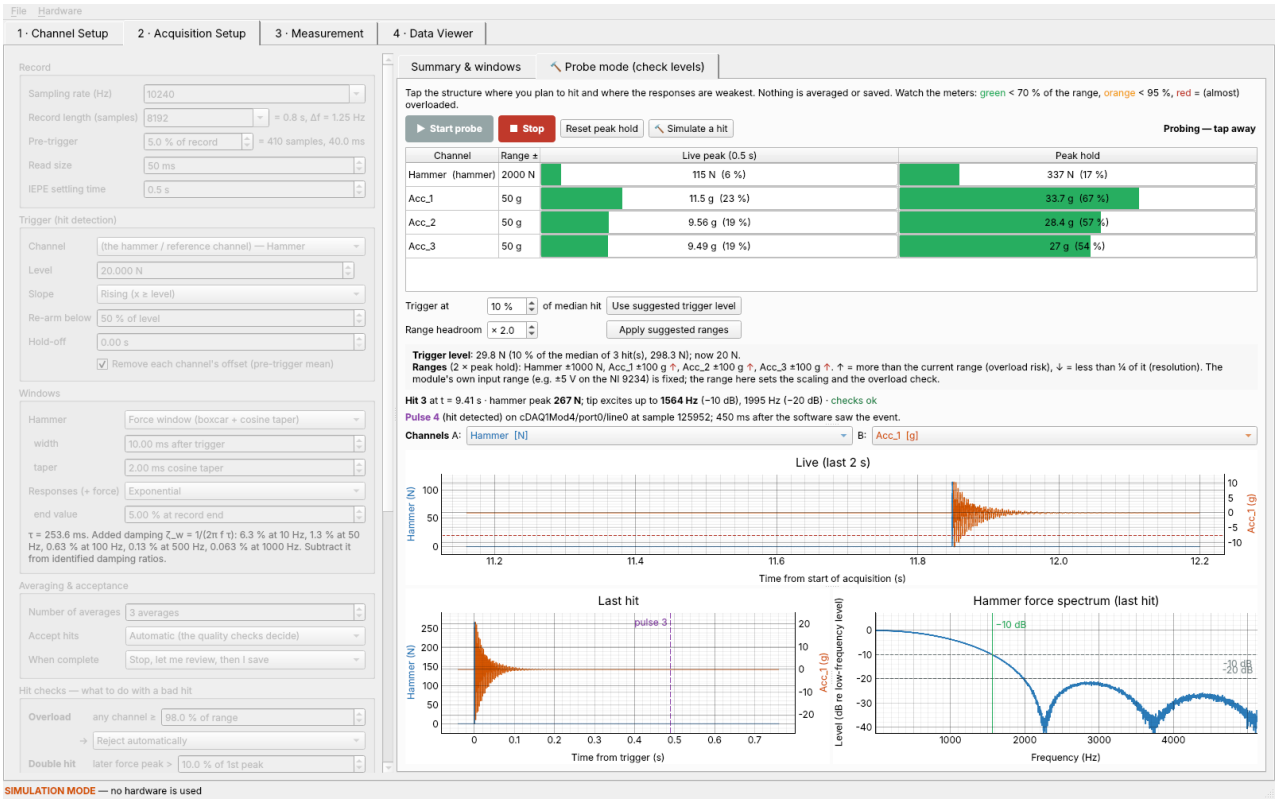


Figure 11: Probe mode: level meters with peak hold, suggestions, the last hit's checks and hammer bandwidth, the live signals, the last hit and the hammer force spectrum.

Press **Start probe** and tap the structure where you plan to hit, and also where the responses are weakest.

- **Level meters**: for every channel, the live peak (last 0.5 s) and the peak hold, as a percentage of its range: green below 70 %, orange below 95 %, red at 95 % and above. *Reset peak hold* starts again.
- **The last hit**: its time history (channels A and B) and the **hammer force spectrum** in dB, with the −10 dB and −20 dB lines. The hammer tip excites the structure reliably up to about the −10 dB frequency, which is printed under the meters. If it is below the highest frequency you care about, use a harder tip (or a lighter hammer).
- **Checks**: every probe hit is checked (double hit, overload, ...) and the result is printed.
- **Suggestions**: a trigger level (a chosen percentage of the median hammer peak) and input ranges (a headroom factor times the peak hold, rounded up to 1, 2 or 5 × 10^n). An up arrow (↑) marks a range larger than the present one (overload risk), a down arrow (↓) one less than a quarter of it (poor resolution). *Use suggested trigger level* sets it on the left; *Apply suggested ranges* writes the ranges into the Channel Setup table after asking.
- **Simulate a hit** (simulation mode only) strikes the simulated structure at once.

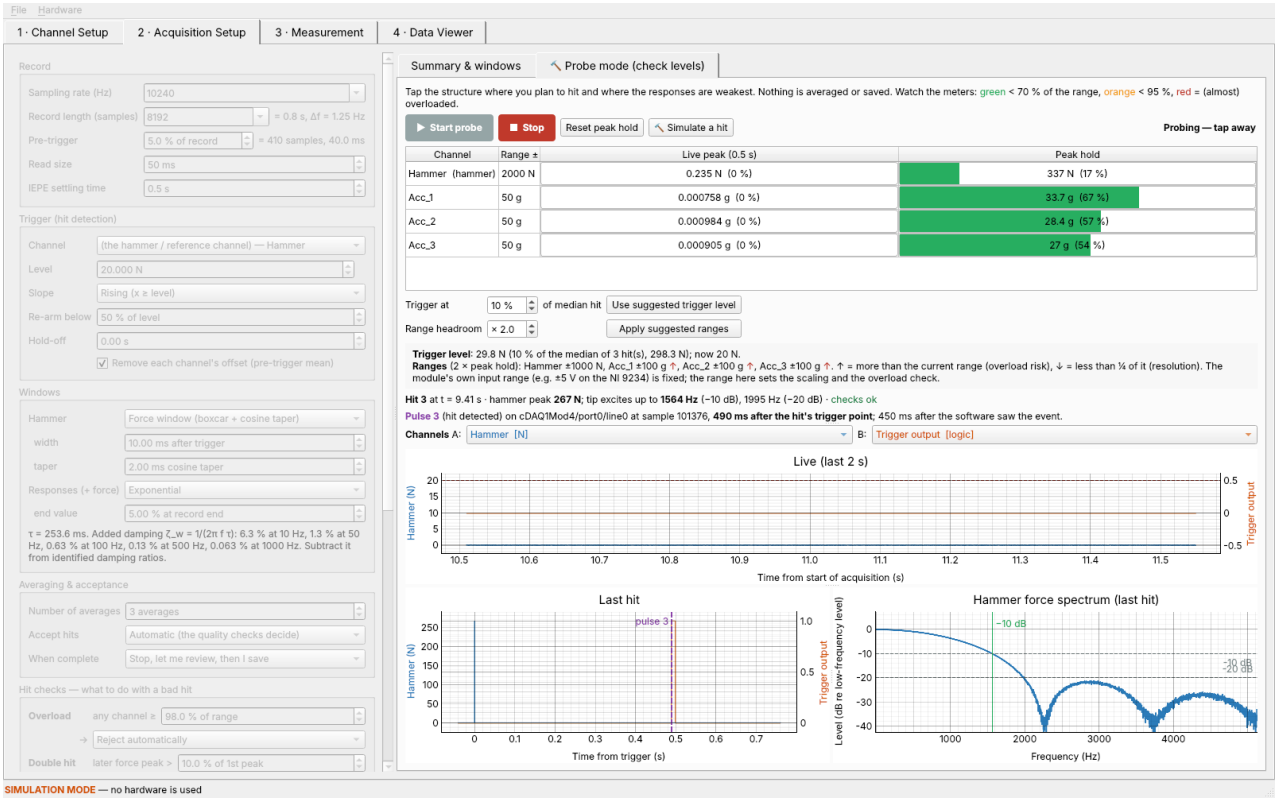


Figure 12: Probe mode with the external trigger: the trigger output is recorded and shown as channel B. The hit is at t = 0; the pulse (purple marker, “pulse 1”) leaves about 450 ms later, i.e. one output queue.

With **Also send pulses in Probe mode** ticked, the pulses are shown too. With hardware timing the trigger output is recorded sample by sample and can be chosen as channel A or B (*Trigger output*, chosen as B by default); every pulse is marked with a purple dashed line, and the panel states its sample index and its delay after the hit. With on-demand timing there is no recorded waveform and only the markers are shown. This is the easiest way to check the camera wiring and the delay to set in the camera before measuring.

8 Measurement

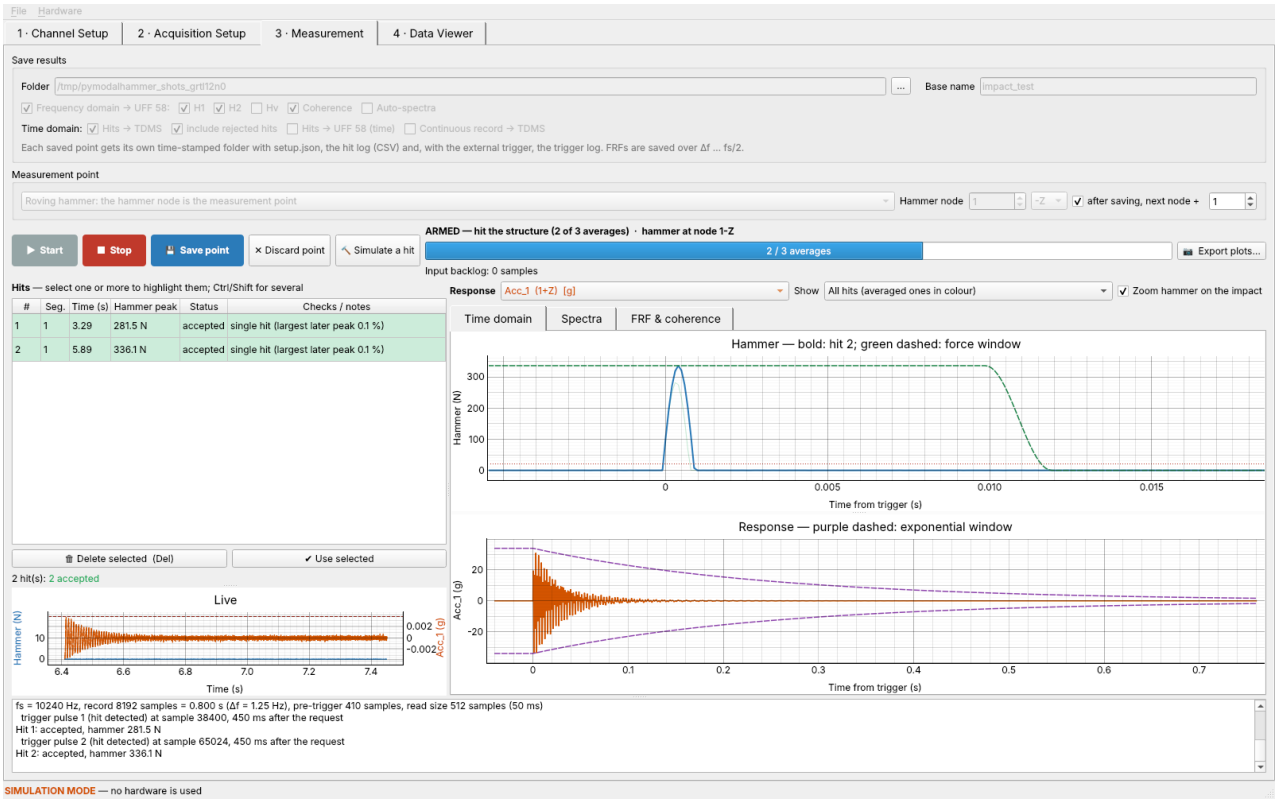


Figure 13: The Measurement tab while a point is being measured: save options and the measurement point at the top, the run controls, the hit list with a live view (left) and the hit plots (right).

8.1 Save options

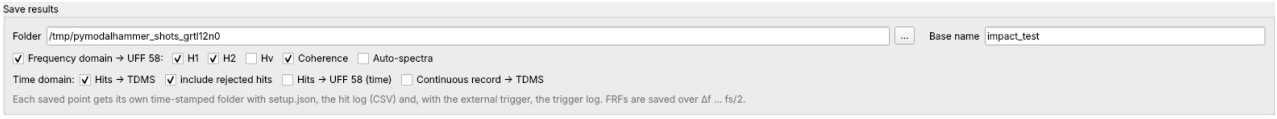


Figure 14: Save options.

- **Folder** and **Base name**: every saved point gets its own sub-folder here (see 10).
- **Frequency domain → UFF 58**: H1, H2, Hv, coherence, and the averaged auto-spectra.
- **Time domain**: *Hits → TDMS* (the averaged hits; optionally also the rejected ones), *Hits → UFF 58 (time)*, and *Continuous record → TDMS* (everything acquired, as a safety copy).
- Always saved: the setup (`setup.json`), the hit log and, with the external trigger, the trigger log.

8.2 Measurement point

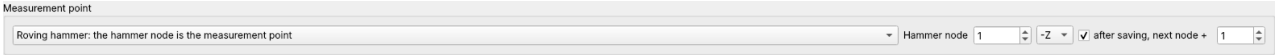


Figure 15: The measurement point.

- *Roving hammer* (default): the hammer is moved from point to point and the accelerometers stay. The **Hammer node** and **direction** given here are the reference node and

direction of the point and are written into the UFF files. With *after saving, next node + n* ticked, the node advances after every saved point.

- *Fixed hammer*: the hammer stays and the sensors are moved; nodes and directions come from the Channel Setup table.

8.3 Running a point

The buttons:

Button	Action
Start / Continue	starts the acquisition (Continue adds hits to a point that has hits already)
Stop	stops the acquisition; the hits are kept
Save point	writes the point's files; enabled when there are averaged hits and none waits for a decision
Discard point	forgets all hits of the point
Simulate a hit	simulation mode only
Export plots…	saves the visible plots as PNG (2× resolution), SVG or PDF

The line next to the buttons tells you what the program is doing: "Settling sensors…" (the IEPE settling time), "ARMED — hit the structure (n of N averages)", "Waiting for your decision", or "Complete". The bar shows the averages collected, and *Input backlog* warns if the computer cannot keep up. The log at the bottom records every hit, every decision, every trigger pulse and every saved file.

If you change channels, record, trigger, windows or the point while a point has hits, **Continue** asks whether to discard them: hits measured with different settings cannot be averaged together.

8.4 The hit list

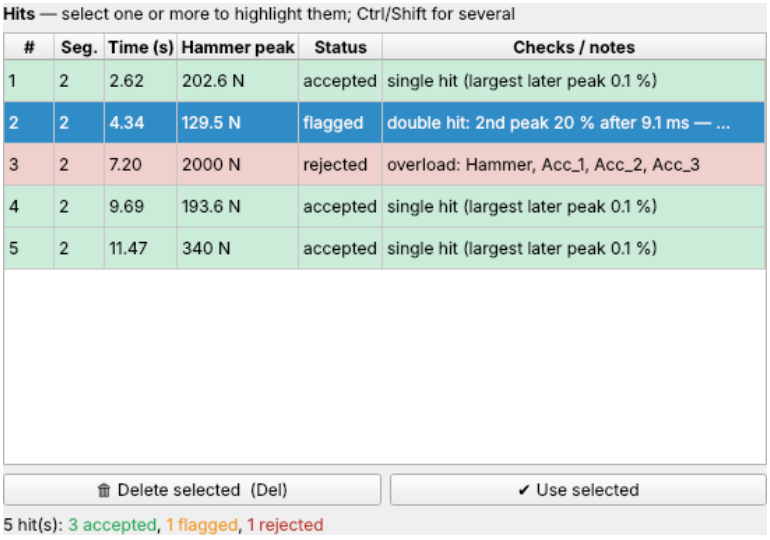


Figure 16: The hit list after the point above: three accepted hits, the double hit that was accepted (flagged, orange) and an overloaded hit rejected automatically (red). Hover over a row for the result of every check.

Colour	Status	Meaning
green	accepted	averaged
orange	flagged	averaged, but a check failed (or you accepted a failed hit)
red	rejected	not averaged (by a check or by you)
blue	pending	waits for your decision

Select one or more rows (Ctrl/Shift-click) to highlight those hits in every plot.

- **Delete selected** (Del key) deletes the hits for good: they leave the list and the average and are never saved.
- **Use selected** puts rejected hits into the average.

The FRF updates immediately.

8.5 Manual review and “Ask me”

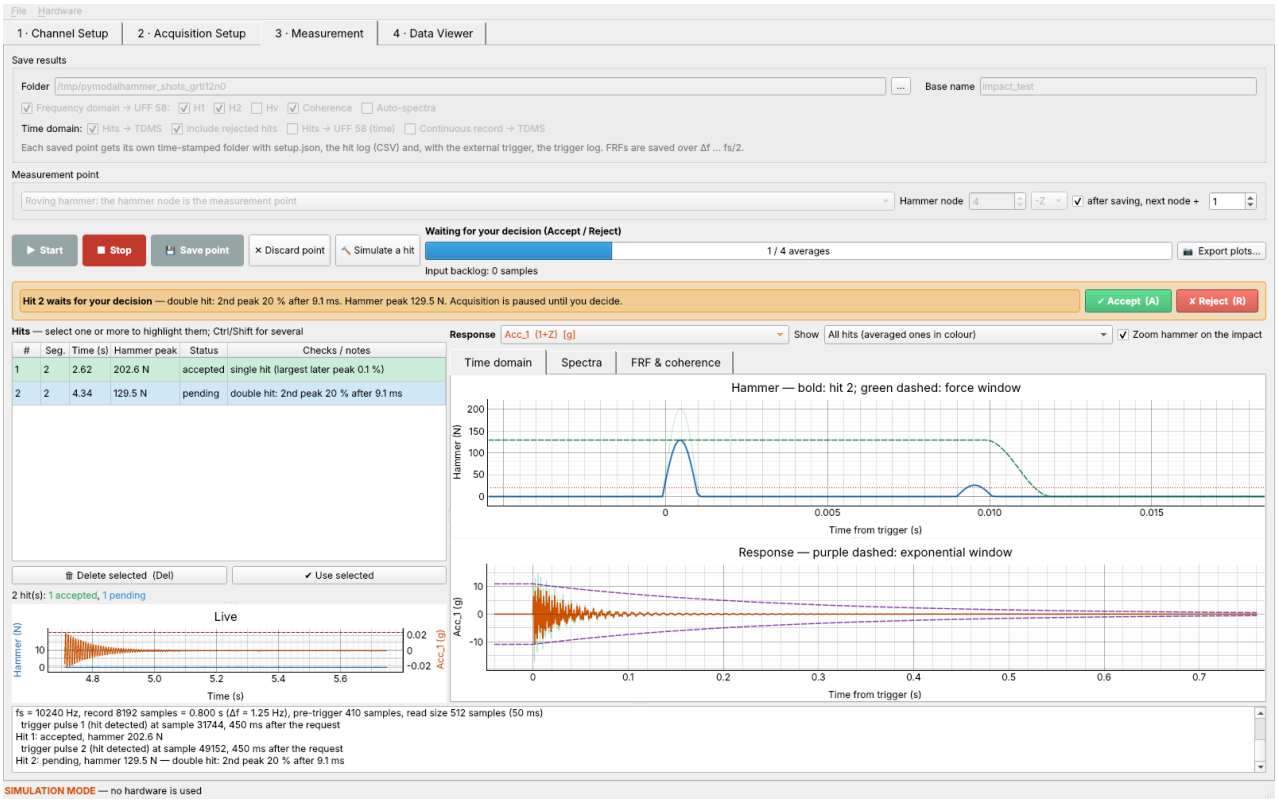


Figure 17: A hit waiting for a decision: a double hit, with the action “Ask me”. The hammer plot (zoomed on the impact) shows the second, smaller impact a few milliseconds after the first, still inside the force window.

When a hit needs your decision, an orange bar names the hit and the reason. Look at the plots and press **Accept** (key A) or **Reject** (key R). Acquisition pauses (the trigger is disarmed) until you decide.

8.6 Time domain, spectra and FRF

Above the plots: the **response** channel shown, **Show** (all hits with the averaged ones in colour, averaged hits only, or only the selected / latest hit), and **Zoom hammer on the impact**.

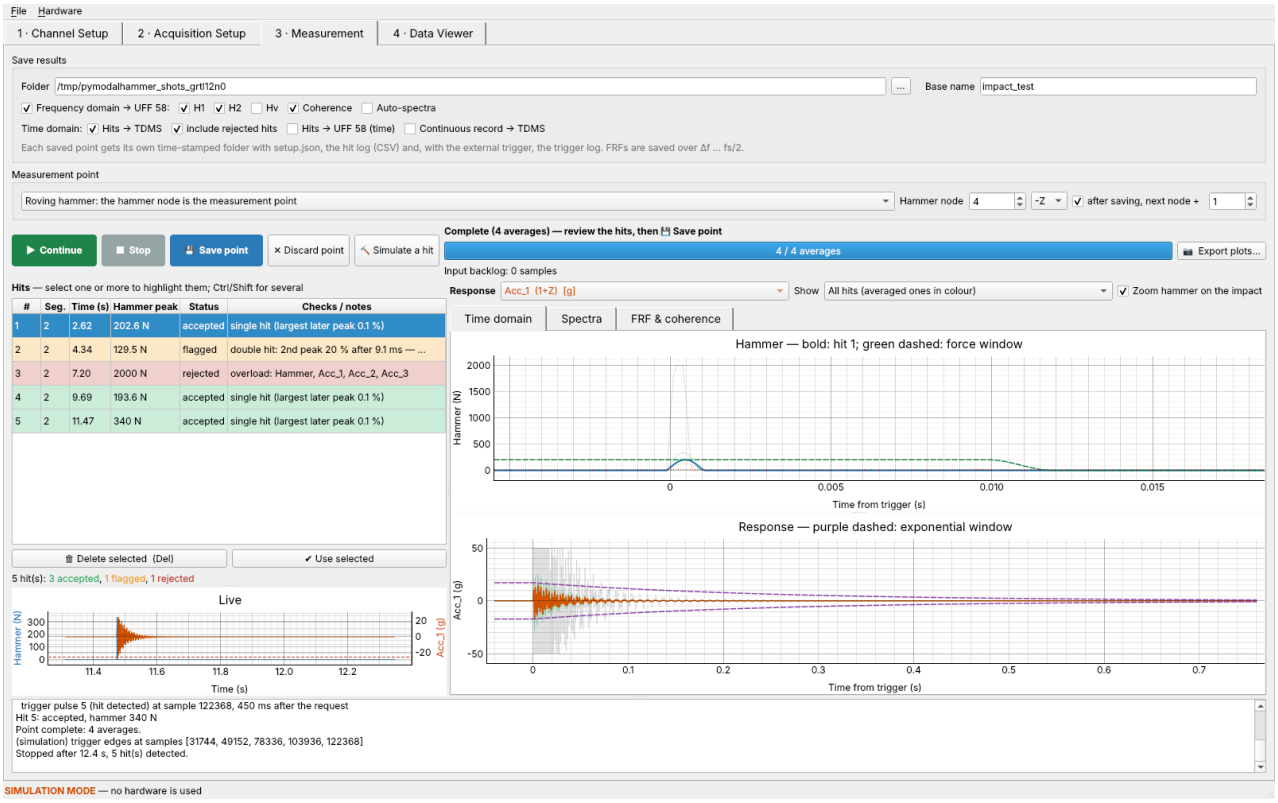


Figure 18: Time domain: the hammer (top, zoomed on the impact, force window dashed green, trigger level dotted red) and the response (bottom, exponential window dashed purple). The selected hit is bold; the others are faint in their status colour.

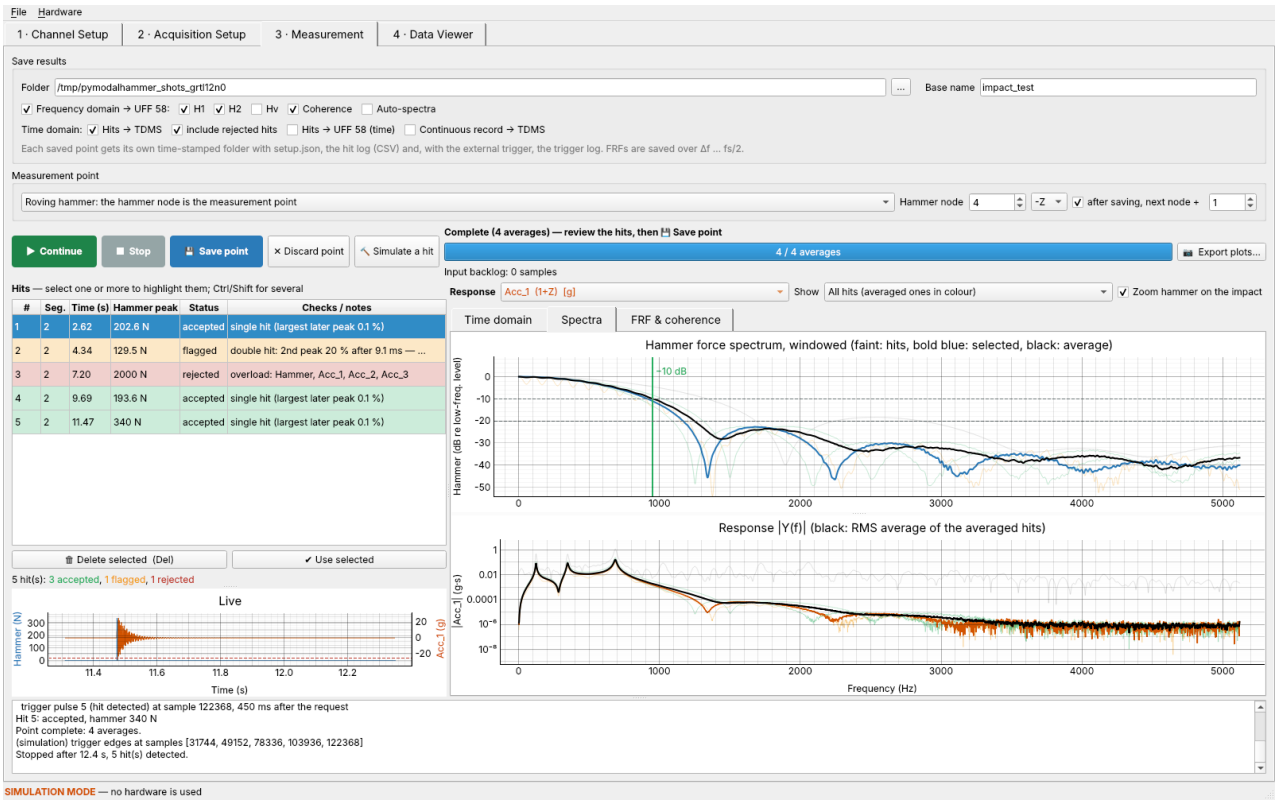


Figure 19: Spectra: the hammer force spectrum in dB (−10 dB frequency marked) and the response spectrum, per hit and averaged (black).

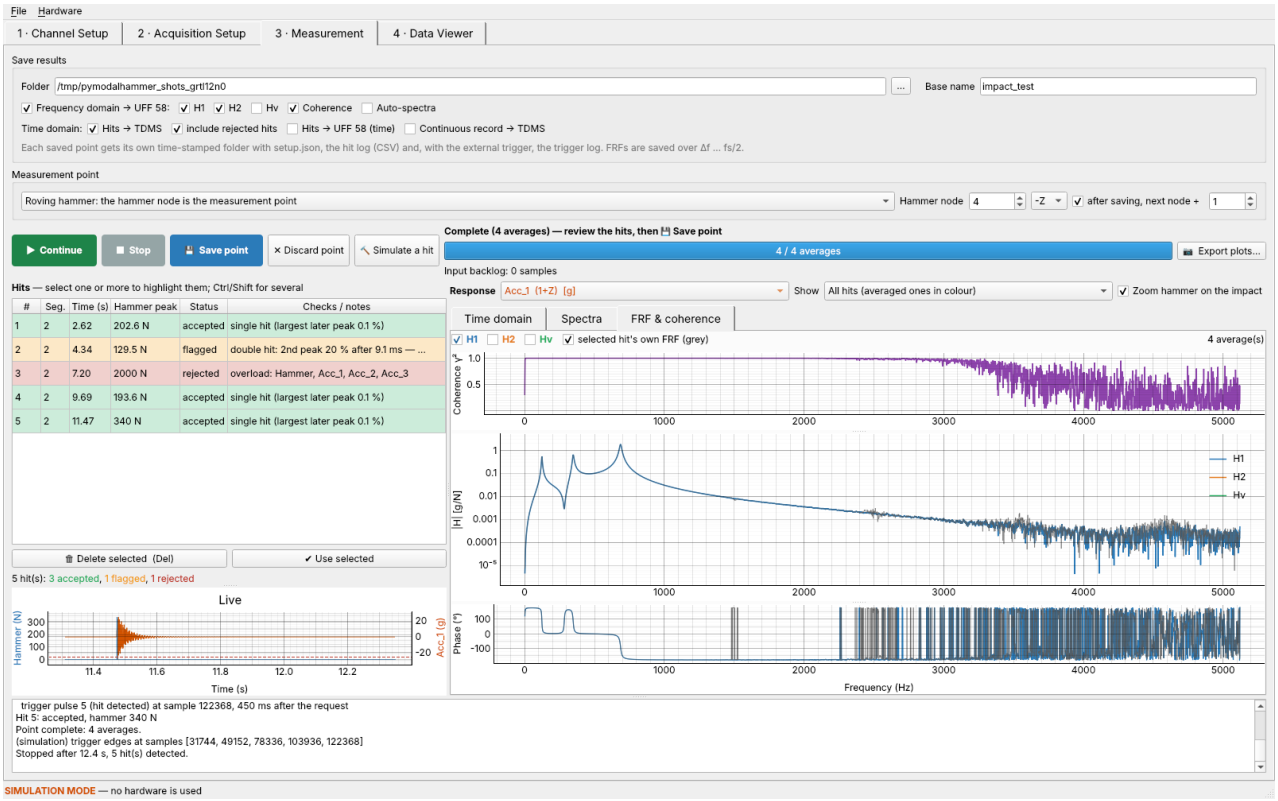


Figure 20: FRF and coherence. The grey curve is the selected hit's own FRF, useful to spot a hit that does not fit. The coherence falls above about 3 kHz because the simulated hammer tip hardly excites the structure there (compare the hammer spectrum).

Reading the coherence: close to 1 where the hits agree and there is enough excitation; low at anti-resonances (little response) and above the hammer tip's bandwidth (little force). A dip at a resonance suggests leakage (use a longer record or the exponential window) or a non-linearity.

8.7 Roving-hammer workflow, step by step

1. Mount the accelerometers; enter their nodes and directions in the Channel Setup table.
2. Probe the structure; set the trigger level and ranges from the suggestions.
3. On the Measurement tab choose *Roving hammer*, the first node and the hammer direction, and tick *after saving, next node + 1*.
4. Choose *When complete*: *Stop, let me review* for careful work, or *Save, then re-arm* for speed.
5. Start, hit the first point N times; review; save. The node advances.
6. Move to the next point and continue until the grid is done.
7. In the Data Viewer, compare the points (each curve is one hammer position).

8.8 Time data only (hammer channel alone)

For impact experiments where only the force time history matters (no response sensors, no modal analysis), enable just the hammer channel and tick *Reference* on it. The measurement then runs as usual, with these differences:

- Record, trigger, pre-trigger, read size, settling and hold-off work exactly as before.

- The hit checks still apply (overload, double hit, force level); the decay check is skipped, since there is no response to decay.
- *Averages* is the number of hits to keep for a point; *When complete* works as before.
- The external trigger output works in every mode (start, every detected hit, every accepted hit, first hit of the point, point complete, frame sync), and every pulse is logged.
- No FRFs or auto-spectra are computed. The response plots and the *FRF & coherence* view are hidden; the hammer's time data and force spectrum are still shown.
- A saved point contains `setup.json` (with `"time_only": true`), the hit log, the trigger log, the hits in TDMS and, if ticked, the hits as UFF 58 time records, and the continuous record if chosen. The frequency-domain UFF files are not written, even if ticked.
- The Data Viewer opens such points on the *Time data* tab; the FRF tab stays empty.

9 Data Viewer

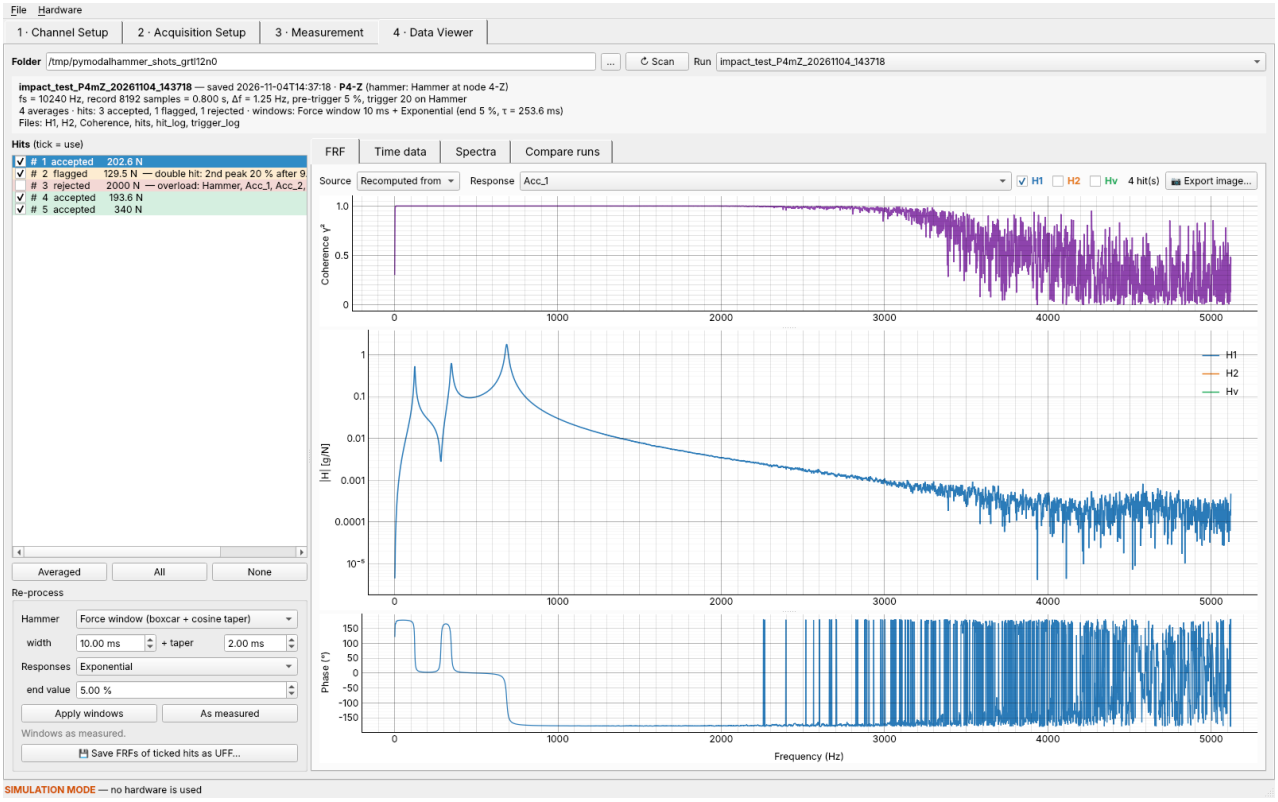


Figure 21: Data Viewer, FRF page: the hits of the chosen run (left, ticked = used), the re-processing box, and the FRF recomputed from the ticked hits.

Enter a run folder or a folder of runs (press *Scan*), then choose a run. The box under the folder summarises the run: point, sampling, record, averages, hits, windows and the files present. A point you have just saved opens here automatically (during a roving-hammer run that keeps acquiring, it opens when you go to this tab).

9.1 Hits and re-processing

The list on the left shows every saved hit with its status. **Ticked hits are used** for the recomputed FRF and the spectra; *Averaged*, *All* and *None* change the ticks. The highlighted row is the hit drawn bold in the time and spectrum plots.

The **Re-process** box applies other windows to the ticked hits (*Apply windows*, back with *As measured*). *Save FRFs of ticked hits as UFF…* writes H1, H2, Hv and coherence of the ticked hits, with these windows, to a folder you choose. The run itself is never modified.

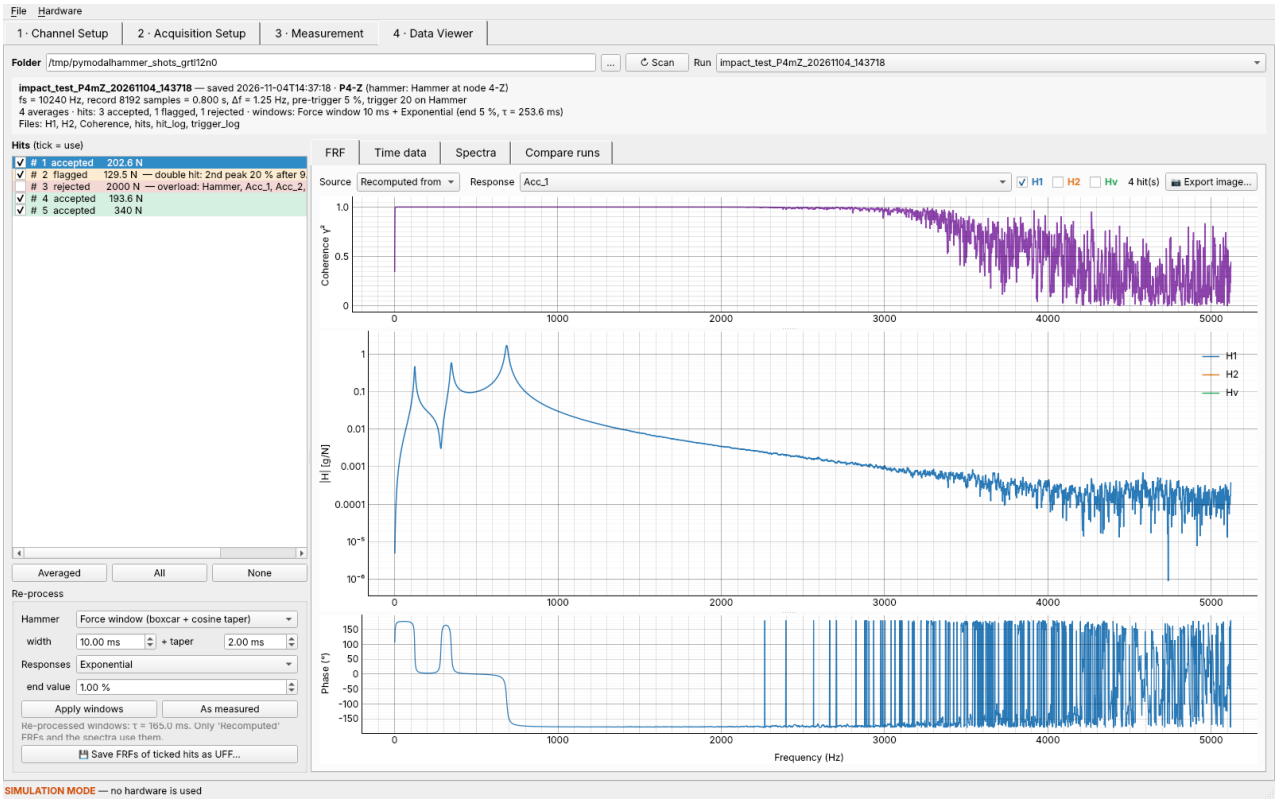


Figure 22: Re-processing with a stronger exponential window (end value 1 % instead of 5 %): the resonance peaks become lower and wider (added damping).

9.2 FRF, time data and spectra

- **FRF**: source *Recomputed from the ticked hits* or *Stored UFF files* (exactly as saved); response channel; H1/H2/Hv; coherence, magnitude and phase.
- **Time data**: two channels (A, B) of all ticked hits (faint) and the selected hit (bold), with the windows and optionally the external-trigger output.
- **Spectra**: the Fourier transform |X(f)| of one channel for every ticked hit, the selected one and their RMS average; windowed or not; logarithmic or dB.

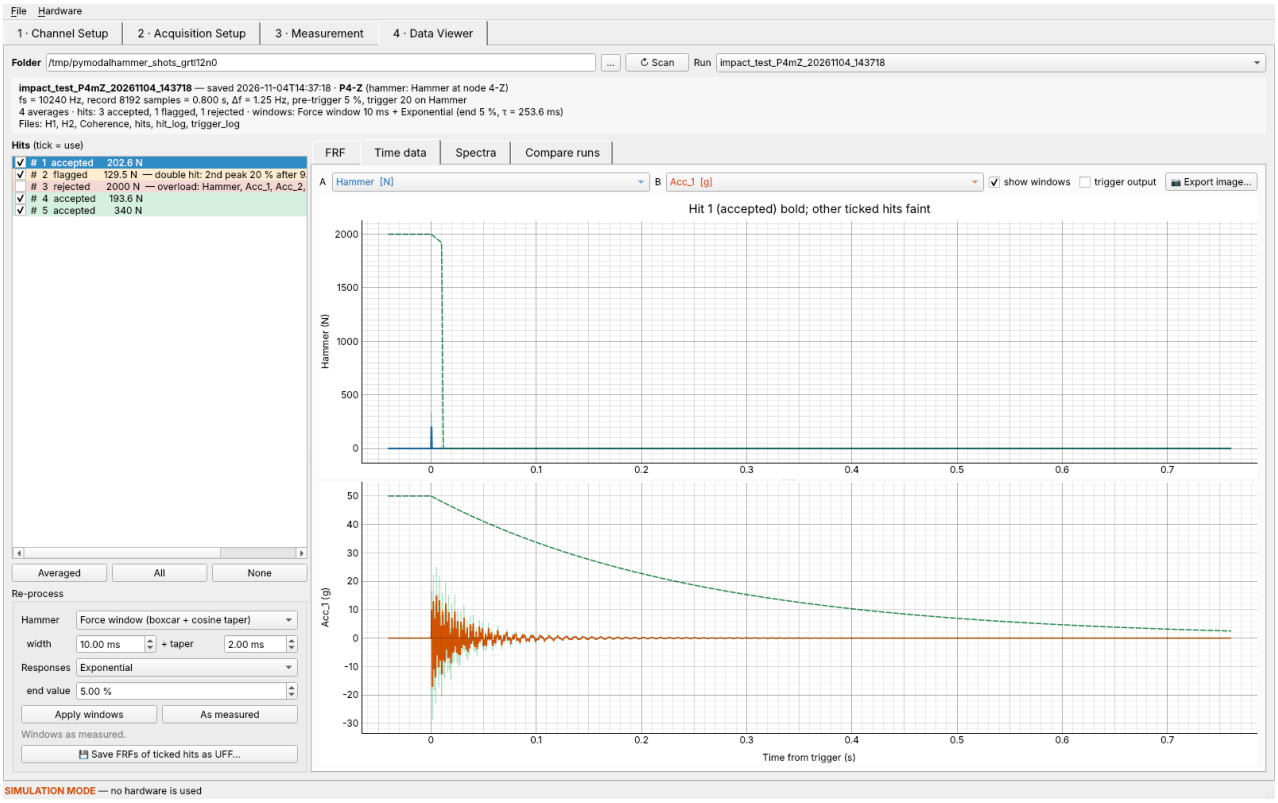


Figure 23: Data Viewer, time data of the hits of a point.

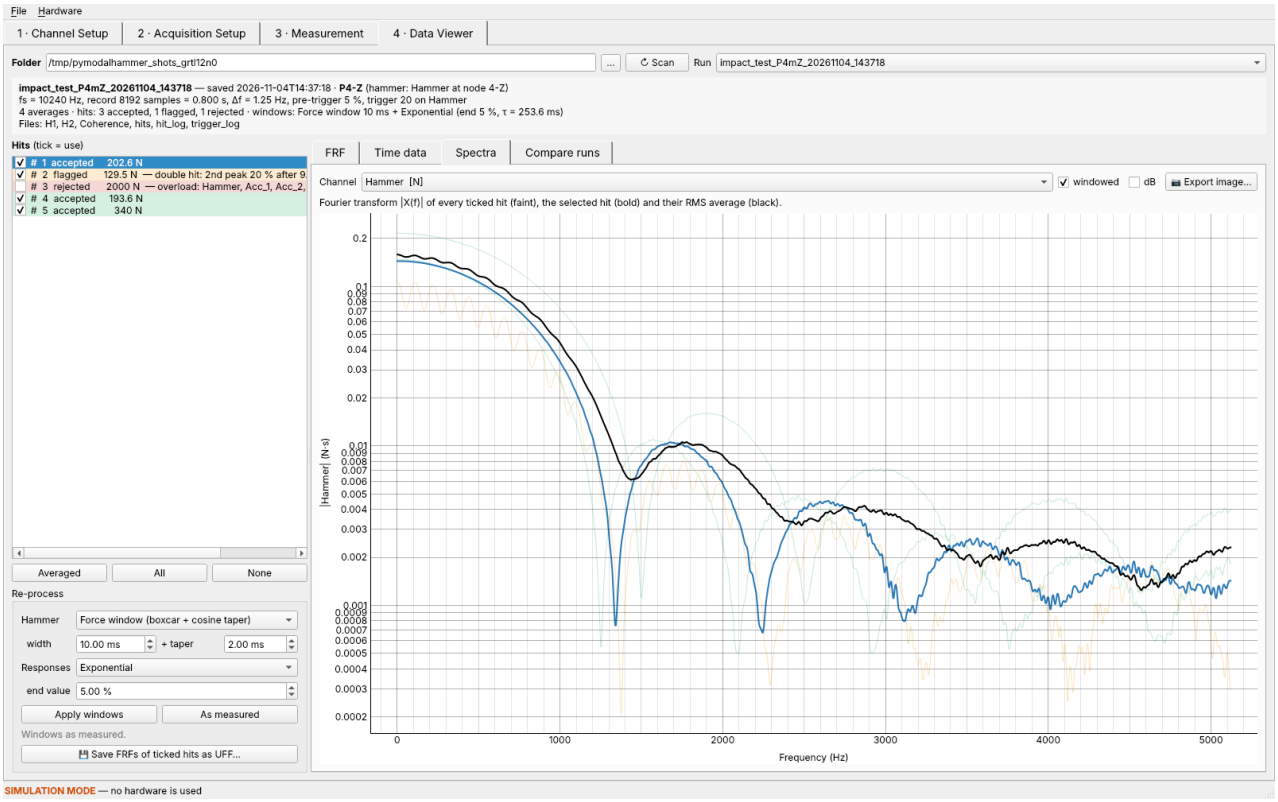


Figure 24: Data Viewer, spectra of the hammer for every hit and their average.

9.3 Compare runs

Tick the runs to compare on the left (*Add runs from folder…* adds runs from elsewhere), choose what to show and the channel. Channels are matched by **name**; each run uses its own hammer reference, so in a roving-hammer test every curve is one hammer position. The legend shows the point of each run.

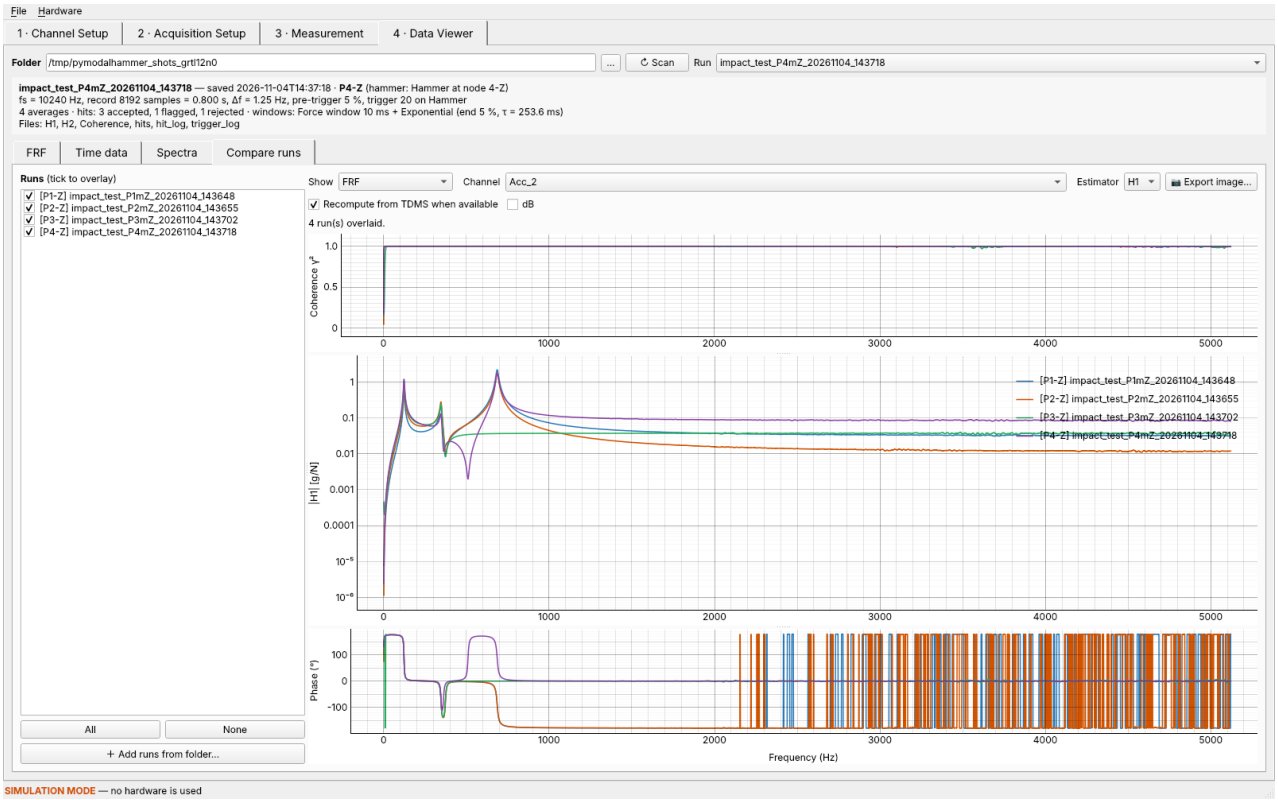


Figure 25: Comparing the FRF of one accelerometer for four hammer points: the mode shapes make the curves differ.

- **FRF**: H1/H2/Hv with coherence and phase. *Recompute from TDMS when available* uses each run's averaged hits; otherwise the stored UFF files. *dB* shows the magnitude in dB.
- **Amplitude spectrum**: the RMS-averaged |X(f)| of the channel.
- **Time data**: the *first averaged hit*, *every averaged hit*, a given *hit number*, or the *mean of the averaged hits* of each run, all aligned at the trigger point. ÷ *hammer peak* divides each record by its hit's peak force, so hits of different strength can be compared; *windowed* applies each run's windows; *zoom on impact* shows the first 20 ms.

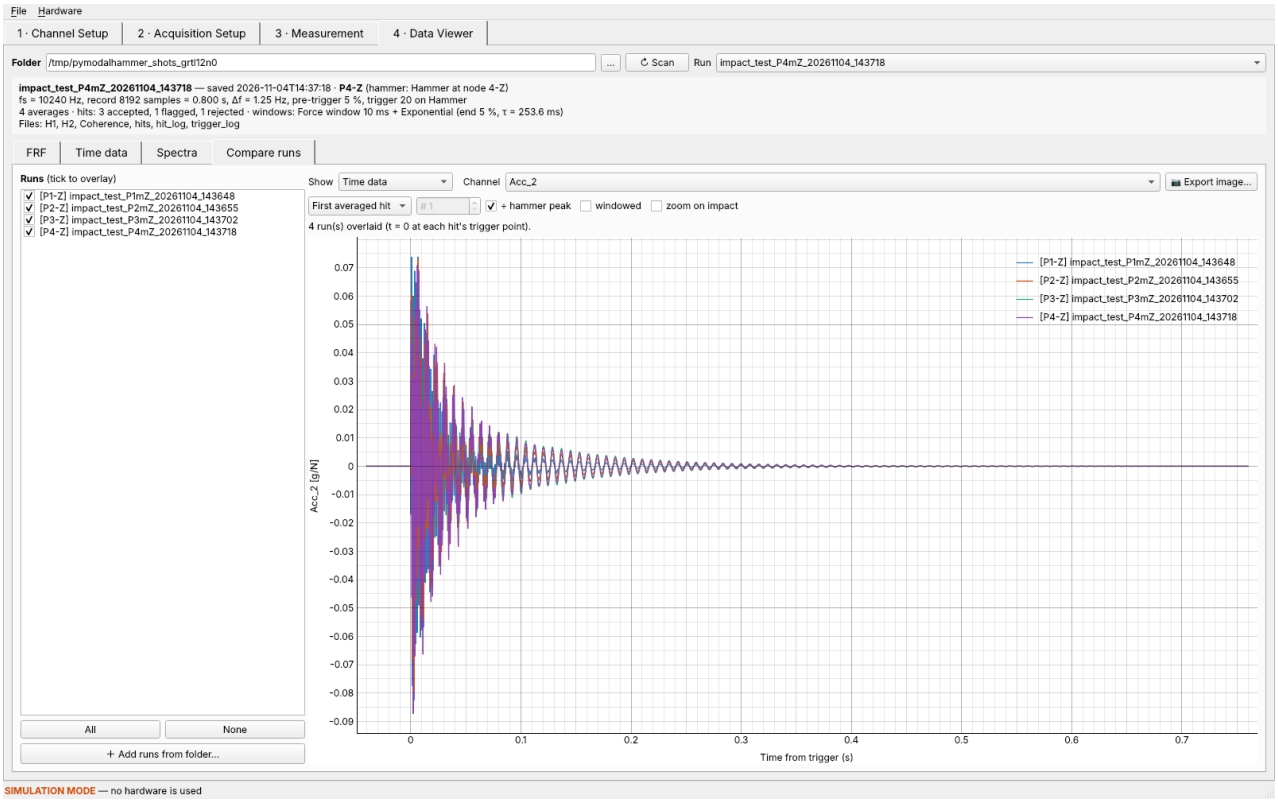


Figure 26: Comparing time data: the first averaged hit of each point, divided by its hammer peak.

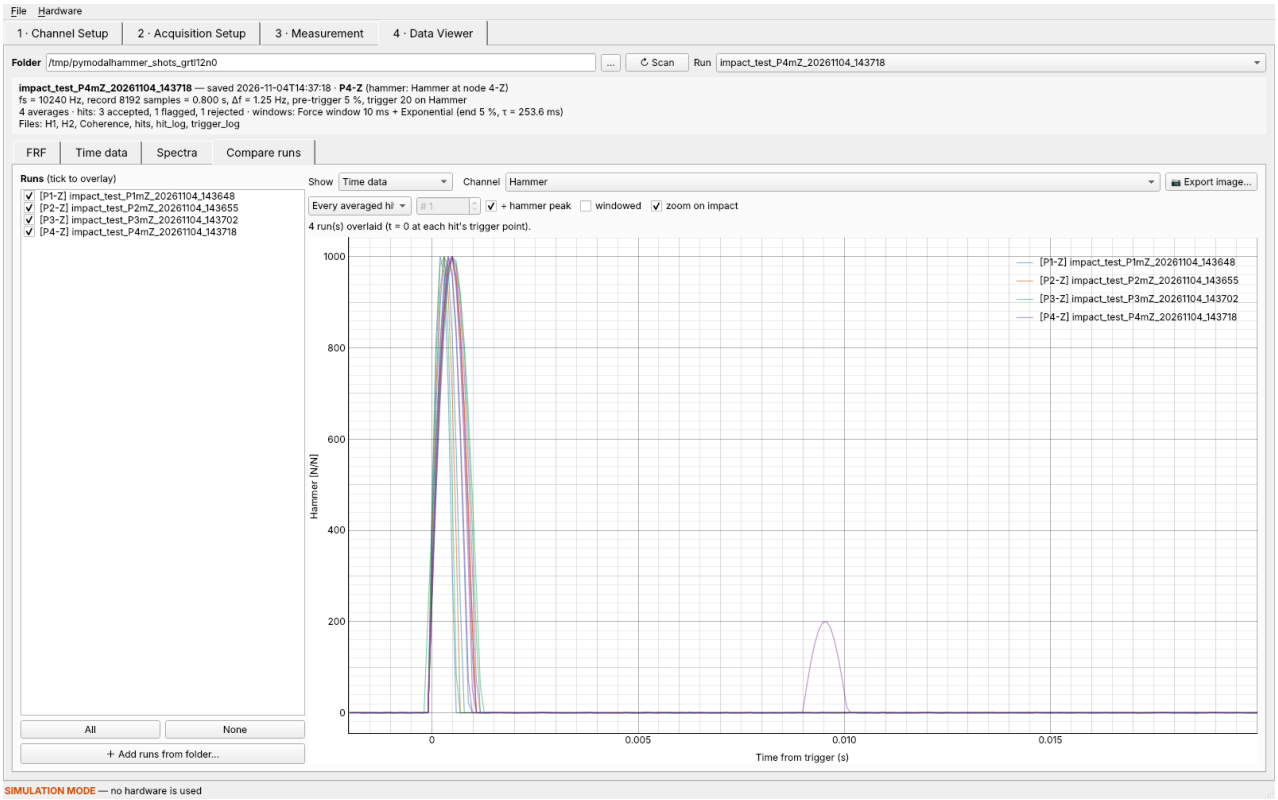


Figure 27: Comparing the hammer pulses of every averaged hit of four points (zoom on impact): a quick check that the hits were alike.

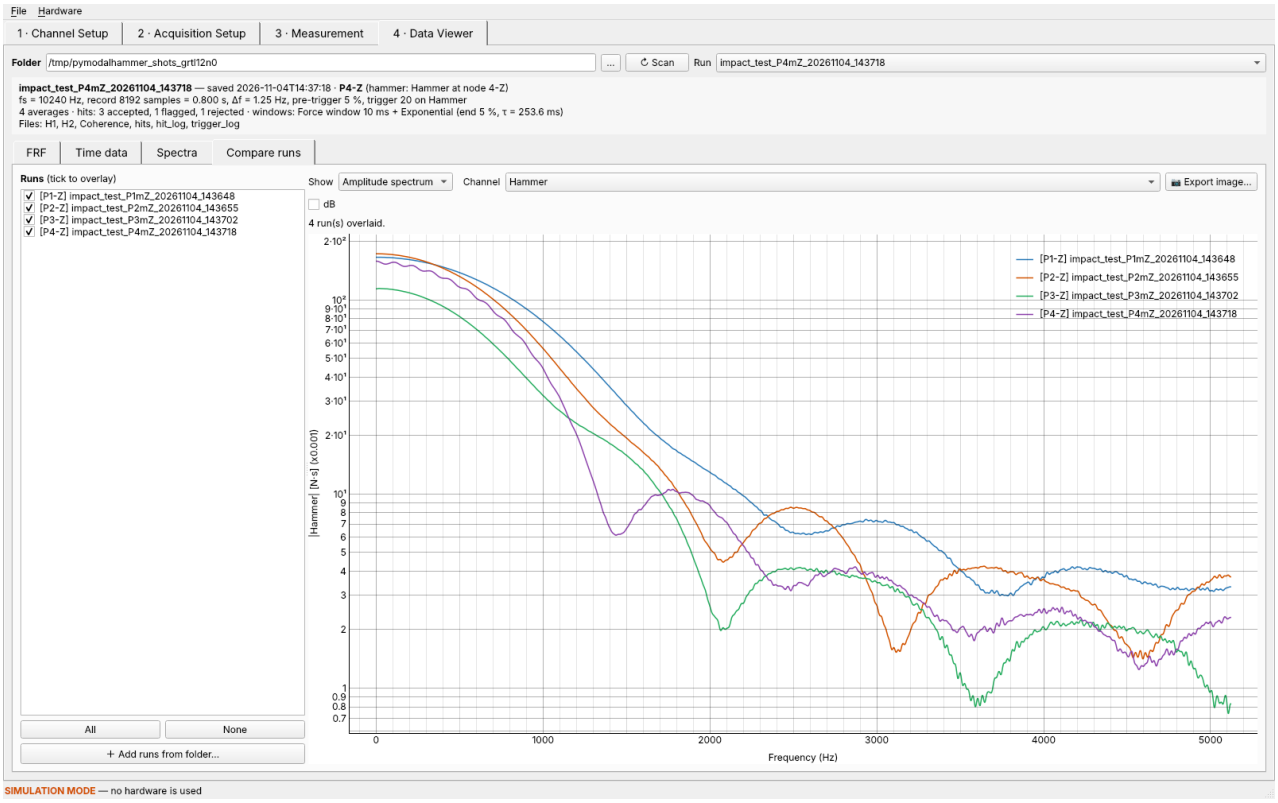


Figure 28: Comparing the hammer amplitude spectra of four points.

Every view has an *Export image...* button (PNG at 2× resolution, SVG or PDF).

10 Saved files

Each saved point goes to its own folder:

```
roving hammer:   <folder>/<base>_P<node><direction>_<date>_<time>/
                 e.g.  impact_test_P3mZ_20261104_142533/
fixed hammer:    <folder>/<base>_fixed_<date>_<time>/
```

(`p` and `m` stand for + and − in the direction; `_2`, `_3` … are added if two points are saved within the same second.)

File	Contents
`setup.json`	every setting, plus a `run` block: actual sampling rate, number of averages, hit counts, τ of the exponential window
`<base>_H1.uff` (H2, Hv)	UFF dataset 58, function type 4 (FRF), complex, one record per response, from Δf to fs/2
`<base>_Coherence.uff`	UFF 58, function type 6
`<base>_AutoSpectra.uff`	optional: averaged auto-spectra of the hammer and every response, function type 2
`<base>_time.uff`	optional: the averaged hits as time responses, function type 1, t = 0 at the trigger point
`<base>_hits.tdms`	the averaged hits (optionally also the rejected ones), one group per hit
`<base>_hit_log.csv`	one line per hit
`<base>_trigger_log.csv`	one line per external-trigger pulse
`<base>_stream_<time>.tdms`	optional, in the main folder: the complete continuous record

10.1 UFF 58 records

Each record carries the response name, node and direction, and the reference (hammer) name, node and direction (direction codes ±1/2/3 for ±X/Y/Z), the units of response and reference, and an ID line such as `H1 Acc_2 5+Z / Hammer 3-Z`. The frequency axis starts at Δf (DC is left out).

10.2 TDMS hit file

- File properties: base name, sampling rate, record length, pre-trigger samples, trigger level and slope, windows, averages, point node/direction/mode and the external-trigger settings.
- One group per hit, `Hit_001`, `Hit_002` …, with these properties:

Property	Meaning
status, included, reason	as in the hit list; included is true for averaged hits
checks	the result of every check
trigger_sample, trigger_time_s	the trigger point, from the start of the segment
segment, wall_time	acquisition segment (each Start or Continue), computer clock time
pre_samples	pre-trigger length in samples
trigger_pulses	sample indices of the external-trigger pulses tied to the hit
peak_<channel>	largest absolute value of each channel

- One channel per sensor, plus `Trigger` with the hardware-timed trigger output. Every channel carries `wf_increment` (1/fs) and `wf_start_offset` (minus the pre-trigger time, so that t = 0 is the trigger point), so programs that read TDMS waveforms get the right time axis. Further properties: unit, node, direction, sensor type, sensitivity, physical channel and range.

10.3 CSV logs

The hit log, `<base>_hit_log.csv`, has one line per hit:

Column	Meaning
hit, segment	hit number within the point, acquisition segment
status, included	accepted / flagged / rejected; 1 if averaged
trigger_sample	sample index of the trigger point from the start of the segment
trigger_time_s, wall_time	the same in seconds; the computer's clock time
peak_<channel>_<unit>	one column per channel: the largest absolute value in the record
reason	why the hit was flagged or rejected, and your decisions
trigger_pulse_samples	sample indices of the external-trigger pulses tied to this hit

The trigger log, `<base>_trigger_log.csv`, has one line per external-trigger pulse:

Column	Meaning
pulse, segment	pulse number, acquisition segment
reason, hit	why it was sent (start, hit detected, accepted hit, point complete …), which hit
sample_index, time_s	when the pulse started, from the start of the segment
sample_exact	1 for hardware-timed pulses (exact), 0 for on-demand pulses (approximate)
requested_at_sample	the sample at which the program asked for the pulse
delay_after_request_ms	the delay between the request and the pulse (≈ the output queue)

10.4 Reading the files in Python

```
from nptdms import TdmsFile
import pyuff

f = TdmsFile.read("impact_test_hits.tdms")
for g in f.groups():
    if g.properties["included"]:
        acc = g["Acc_1"]
        t = acc.time_track()              # seconds; 0 = trigger point
        print(g.name, g.properties["status"], acc[:].max())

for rec in pyuff.UFF("impact_test_H1.uff").read_sets():
    print(rec["id1"], rec["x"][:3], rec["data"][:3])
```

11 How it works

11.1 Hit detection

The inputs are read in pieces (*read size*) into a ring buffer that always holds more than one record. In every new piece the trigger channel is searched for the first sample that crosses the level, after the signal has dropped below the re-arm level. A crossing at sample s starts the record [s − pre, s − pre + N). When its last sample has arrived, the record is copied, the pre-trigger mean is subtracted (if chosen) and the checks run. The next search starts after the end of the record (plus the hold-off).

11.2 Spectra and FRF estimators

Each hit's spectra are the windowed Fourier transforms $X(f) = \mathrm{DFT}(x \cdot w)/f_s$, in EU·s. With X the hammer and Y a response, averaged over the hits used:

$$G_{xx} = \langle |X|^2 \rangle, \quad G_{yy} = \langle |Y|^2 \rangle, \quad G_{xy} = \langle X^* Y \rangle$$

$$H_1 = \frac{G_{xy}}{G_{xx}}, \qquad H_2 = \frac{G_{yy}}{G_{xy}^*}, \qquad H_v = \frac{G_{xy}}{G_{xx} - \lambda_{\mathrm{min}}}, \qquad \gamma^2 = \frac{|G_{xy}|^2}{G_{xx}\, G_{yy}}$$

where λ_{min} is the smaller eigenvalue of $\begin{bmatrix} G_{xx} & G_{xy}^* \\ G_{xy} & G_{yy} \end{bmatrix}$. H1 suits noise on the response, H2 noise on the force; Hv lies between them. For an impact test H1 is the usual choice.

11.3 Checks in detail

- **Overload**: max |x| of each channel compared with the given percentage of its range.
- **Double hit**: from the hammer's peak, the end of the pulse is the first sample below 5 % of the peak; 0.3 ms later the remaining record is searched for |force| above the given percentage.
- **Not decayed**: $\sqrt{2}$ RMS of the last 5 % of each response record, as a percentage of that response's peak.

12 Hardware notes

- **Sampling rates.** Delta-sigma modules (NI 9234, 9250, 9260 …) accept only certain rates (51.2 kHz / n); the rate really used is read back, logged and saved.
- **IEPE.** The NI 9234 and 9250 supply 2 mA only. Use AC coupling for IEPE sensors and allow a few seconds of settling.
- **Filter delay.** Delta-sigma inputs delay every sample by the module's filter group delay (≈ 38 samples on the NI 9234). It does not affect the FRFs (all inputs share it), but an external-trigger pulse is practically undelayed, so camera frame f matches input sample f + delay. For sub-sample certainty, also record the trigger line on a spare input.
- **Analog trigger outputs on the NI 9260** (delta-sigma): use hardware timing, pulses of ≥ 10 samples, and a level within ±4.24 V. Its reconstruction filter delays and rounds the pulse; a digital line (NI 9401) gives clean edges.
- **Digital lines on the NI 9401** are set as inputs or outputs in groups of four; its outputs are 5 V TTL.
- **USB.** Keep the read size ≥ 20 ms. The *Input backlog* indicator warns if the computer falls behind.

13 Troubleshooting

Symptom / message	Cause and remedy
No hits are detected	Trigger level too high, or wrong channel or slope. Probe the structure and use the suggested level. Wait for `ARMED`.
Hits trigger by themselves	Level too close to the noise: raise it, or raise *Re-arm below* and the hold-off.
Many double hits	Let the hammer rebound freely; use a lighter hammer or a harder tip; strike with a flick of the wrist.
Overloads	Hit more softly or raise the channel's range (Probe mode suggests one).
Coherence low at high frequencies	The tip does not excite them: use a harder tip; check the hammer spectrum's −10 dB frequency.
Peaks too wide / damping too high	The exponential window adds damping: weaken it, lengthen the record, or subtract the added damping.
"Task contains physical channels … require you to specify the Sample Clock rate" (−201087)	The output module (e.g. NI 9260) cannot be switched on demand: use hardware timing.
"Requested sample clock source is invalid" (−200414)	Only with *Clocked by the AI sample clock*; the program falls back by itself. Choose *Own clock* to avoid the message.
"Onboard device memory underflow" (−200621)	The trigger output ran dry: raise *Output queue* (1000–2000 ms), close other programs, or use a DO line on demand.
"PC is falling behind"	Increase the read size, close other programs, avoid very high sampling rates over USB.
Window too large for the screen	Every tab scrolls; drag the dividers between panels to give space where you need it.

14 Testing and development

The folder `tests` holds automated tests that run without hardware:

```
# engine, checks, averaging, trigger output, files
python tests/test_core.py
# NI back end against a mock driver with the limits of the NI 9260
python tests/test_ni_mock.py
# the whole GUI workflow in simulation (screenshots go to out/)
QT_QPA_PLATFORM=offscreen python tests/gui_smoke.py out/
# roving hammer with auto-advance
QT_QPA_PLATFORM=offscreen python tests/gui_autoadvance.py
# probe mode shows the trigger pulse
QT_QPA_PLATFORM=offscreen python tests/gui_probe_pulse.py
# hammer channel alone: time data only, with trigger pulses
QT_QPA_PLATFORM=offscreen python tests/gui_time_only.py
```

The tests use their own settings, so they never change your saved setup. The screenshots in

this manual are made by `tools/make_screenshots.py` from a simulated session.

To build this manual as a PDF: open it in Emacs and press `C-c C-e l p`, or run `docs/export_pdf.sh`. It uses XeLaTeX and the DejaVu fonts (both part of common Linux and TeX Live installations).

15 Reporting problems

Please include: the version (in the window title), the operating system, the NI hardware (chassis and modules), the setup file (*File → Save setup as…*), the full text of any error message, and the lines of the log at the bottom of the Measurement tab. For questions about results, the point's folder (or its `setup.json` and `hit_log.csv`) helps most.

16 Licence


```
pymodalhammer — impact-hammer modal testing with NI-DAQmx
Copyright (C) 2026  The pymodalhammer authors

This program is free software: you can redistribute it and/or
modify it under the terms of the GNU General Public License as
published by the Free Software Foundation, either version 3 of
the License, or (at your option) any later version.

This program is distributed in the hope that it will be useful,
but WITHOUT ANY WARRANTY; without even the implied warranty of
MERCHANTABILITY or FITNESS FOR A PARTICULAR PURPOSE. See the GNU
General Public License for more details.

You should have received a copy of the GNU General Public License
along with this program (the file LICENSE). If not, see
<https://www.gnu.org/licenses/>.
```